

Evaluation report

History · narrative · grade band 9-10

Verdict

Stage score	93.7 / 100 (exemplary)
Confidence	92%
Stages scored	10 of 11
Rubric score	92.3 / 100 (exemplary)
Package	bfb360ac-b0da-482d-9781-000b1ff9d6c4
Evaluated	2026-08-03T18:59:12+00:00

Two scores, not one. The stage score asks whether each pipeline stage met its contract; the rubric score asks whether the result is good teaching. They are reported separately because averaging them would produce a number that answers neither question.

Against history

1 comparable run(s) on record; a baseline needs at least 5 before a fall in score can be told apart from run-to-run variation.

Stages

Stage	Score	Conf.	Measured	Unmeasurable
Document intelligence	95.2	50%	3	3
Educational classification	79.5	80%	4	1
Knowledge extraction	88.7	100%	5	0
Teaching planner	100.0	100%	3	0
Lesson generation	100.0	100%	4	0
Activity generation	100.0	100%	6	0
Assessment generation	97.9	86%	6	1
Gap analysis	94.1	100%	5	0
Validation	82.8	100%	4	0
Publishing	97.3	100%	5	0
Learning outcomes	n/a	0%	0	4

Not measurable (9)

These metrics carry no score. Each one states what it would take to measure — a labelled corpus, a classroom trial, response data. Reporting a number here would mean inventing one.

OCR confidence (document-intelligence)

No page needed OCR: every page carried a text layer, so nothing was read from an image and there is no confidence to report. Measuring it would need: a document with at least one scanned page

OCR accuracy (document-intelligence)

Accuracy is a comparison against what the page actually said, and no uploader supplies a transcript of their own scan. Confidence is reported instead, which is the recogniser's certainty, not its correctness — the two come apart exactly when it matters. Measuring it would need: reference transcriptions of the scanned pages to diff against

Extraction fidelity vs. original (document-intelligence)

Fidelity is a comparison, and the original document's ground-truth text is not part of the package. Measuring it would need: a hand-checked transcript of the source PDF to diff the extraction against

Subject / grade classification accuracy (educational-classification)

Accuracy needs labels. There is no corpus of documents with agreed subject and grade-band annotations to score this run against. Measuring it would need: a held-out set of documents labelled by teachers, ideally 100+ across subjects

Item difficulty is calibrated (assessment-generation)

Difficulty is a property of a population's responses. No student has answered these items, so there is no p-value or discrimination index to report. Measuring it would need: response data from a cohort — item-response theory needs roughly 100 students per item

Teacher satisfaction (outcomes)

Satisfaction is a rating teachers give. None have used this package. Measuring it would need: a pilot with teachers rating packages after classroom use — 20+ responses before a mean means anything

Student learning effectiveness (outcomes)

Effectiveness is measured by what students learn, which requires students, a pre-test and a post-test. Measuring it would need: a controlled classroom trial with pre/post assessment against a conventionally-planned unit

Teacher preparation time saved (outcomes)

Time saved is a difference against how long the same teacher would take unaided. That baseline has not been collected. Measuring it would need: timed preparation of the same unit by teachers with and without the package

Conformance to a published curriculum (outcomes)

The package aligns to a board's structure, but no authority has certified the mapping. Reporting a conformance score would state agreement nobody granted. Measuring it would need: an official syllabus document parsed to statement level, and a reviewer's sign-off on the mapping

Recommendations (6)

High

Educational classification — Low-confidence fields flagged for review (33.3)

Derive low_confidence_fields from the confidences map rather than letting the model populate both independently.

Why it matters: A field the classifier doubted reaches teachers unflagged.

Validation — No blocking issues remain (53.3)

Route the blocking issues back to their owning stage for repair.

Why it matters: The package ships with defects its own validator found.

Knowledge extraction — Quotes appear in the chunks they cite (66.1)

Re-run extraction with the citation verifier rejecting spans whose quote is not substring-present in the cited chunk.

Why it matters: A citation that does not check out is worse than no citation — a teacher trusts it.

Assessment generation — Objectives are assessed (90.0)

Drive item generation from the objective list rather than the concept list.

Why it matters: An objective nobody assesses is an objective nobody knows was met.

Medium

Publishing — Every stage reported its cost (90.0)

Emit a timing entry from every stage span, including cheap ones.

Why it matters: Per-stage cost cannot be attributed without it.

Low

Document intelligence — Source metadata captured (85.7)

Record detected_language at parse time.

Why it matters: Provenance cannot be audited without it.

Metric detail

Document intelligence — 95.2

Source metadata captured — 85.7 (measured, confidence 1.00)

6 of 7 source fields recorded; missing detected_language.

- /source — recorded: filename, mime, sha256, size_bytes, page_count, word_count

Text extracted per page — 100.0 (measured, confidence 1.00)

494 words per page across 39 pages. Consistent with a born-digital document.

- /source/word_count — 19,273 words over 39 pages

Chunks carry text — 100.0 (measured, confidence 1.00)

80 of 80 chunks hold text. Empty chunks break every citation that lands on them.

- /chunks — 80 chunks retained

OCR confidence — n/a (not_measurable, confidence 0.00)

No page needed OCR: every page carried a text layer, so nothing was read from an image and there is no confidence to report. Measuring it would need: a document with at least one scanned page

OCR accuracy — n/a (not_measurable, confidence 0.00)

Accuracy is a comparison against what the page actually said, and no uploader supplies a transcript of their own scan. Confidence is reported instead, which is the recogniser's certainty, not its correctness — the two come apart exactly when it matters. Measuring it would need: reference transcriptions of the scanned pages to diff against

Extraction fidelity vs. original — n/a (not_measurable, confidence 0.00)

Fidelity is a comparison, and the original document's ground-truth text is not part of the package. Measuring it would need: a hand-checked transcript of the source PDF to diff the extraction against

Educational classification — 79.5

Classification fields populated — 100.0 (measured, confidence 1.00)

7 of 7 fields classified.

- /classification — subject='History', profile='narrative'

Pedagogy profile is one the pipeline routes on — 100.0 (measured, confidence 1.00)

Profile 'narrative' is a routable profile; every downstream stage branches on it.

- /classification/pedagogy_profile — narrative

Classifier's own confidence — 74.1 (measured, confidence 1.00)

Mean self-reported confidence 0.74 across category, chapter, curriculum_alignment, difficulty, grade_band, language, pedagogy_profile, subject, topic. This is the classifier's own estimate, not an accuracy measurement.

- /classification/confidences — {'subject': 0.95, 'grade_band': 0.6, 'difficulty': 0.7, 'topic': 0.98, 'chapter': 0.6, 'category': 0.9, 'language': 0.99, 'pedagogy_profile': 0.95, 'curriculum_alignment': 0.0}

Low-confidence fields flagged for review — 33.3 (measured, confidence 1.00)

2 field(s) scored below 0.7 without being flagged: chapter, grade_band.

- /classification/low_confidence_fields — ['curriculum_alignment']

Subject / grade classification accuracy — n/a (not_measurable, confidence 0.00)

Accuracy needs labels. There is no corpus of documents with agreed subject and grade-band annotations to score this run against. Measuring it would need: a held-out set of documents labelled by teachers, ideally 100+ across subjects

Knowledge extraction — 88.7

Concepts carry evidence — 100.0 (measured, confidence 1.00)

25 of 25 concepts carry evidence.

Concepts carry a summary — 100.0 (measured, confidence 1.00)

25 of 25 concepts carry summary.

Quotes appear in the chunks they cite — 66.1 (measured, confidence 1.00)

37 of 56 evidence spans quote text that is present verbatim in the chunk they cite. 19 do not.

- /knowledge/concepts/2/evidence — quote not found in c_0008: 'Finding their meeting place closed the next day, they took t'
- /knowledge/concepts/4/evidence — quote not found in c_0011: 'Other decrees included equality before the law, opening\npubl'
- /knowledge/concepts/5/evidence — quote not found in c_0011: 'The declaration was approved by the Assembly and published o'
- /knowledge/concepts/5/evidence — quote not found in c_0011: 'On 9 July, the National Assembly declared itself the Nationa'
- /knowledge/concepts/8/evidence — quote not found in c_0016: 'on 22 September the Convention replaced\nthe monarchy with th'

Objectives carry a Bloom level — 100.0 (measured, confidence 1.00)

10 of 10 objectives are tagged with a Bloom level; 4 distinct level(s) used.

- /knowledge/learning_objectives — levels: ['analyze', 'evaluate', 'remember', 'understand']

Concept graph resolves to real concepts — 100.0 (measured, confidence 1.00)

10 edges over 25 nodes. Every endpoint resolves to an extracted concept.

Teaching planner — 100.0

Every concept is scheduled into a period — 100.0 (measured, confidence 1.00)

25 of 25 concepts are taught in some period.

- /teaching_plan/periods — 6 periods scheduled

Prerequisites are taught before what needs them — 100.0 (measured, confidence 1.00)

6 of 6 prerequisite edges are respected by the period order.

Period plans fit the period — 100.0 (measured, confidence 1.00)

6 of 6 periods allocate no more than the 40-minute period.

Lesson generation — 100.0

Every planned period has classroom content — 100.0 (measured, confidence 1.00)

6 of 6 planned periods have generated content.

- /classroom_content — 6 content blocks for 6 periods

Each period carries the full set of lesson blocks — 100.0 (measured, confidence 1.00)

36 of 36 lesson blocks are populated across 6 periods.

Teacher scripts have enough to work from — 100.0 (measured, confidence 1.00)

Teacher scripts average 1072 words per period against a 250-word working floor. This measures volume, not teaching quality.

- /classroom_content/0/teacher_script — 1050 words

Activity references resolve — 100.0 (measured, confidence 1.00)

6 of 6 activity references point at an activity that exists.

Activity generation — 100.0

Activities vary in kind — 100.0 (measured, confidence 1.00)

5 distinct activity type(s) across 6 activities (debate, field_task, gallery_walk, group_discussion, role_play), against a target of 4.

- /activities — types: ['debate', 'field_task', 'gallery_walk', 'group_discussion', 'role_play']

Activities target extracted concepts — 100.0 (measured, confidence 1.00)

6 of 6 activities reference at least one concept that stage 3 actually extracted.

Activities carry teacher instructions — 100.0 (measured, confidence 1.00)

6 of 6 activities carry teacher_instructions.

Activities state success criteria — 100.0 (measured, confidence 1.00)

6 of 6 activities carry success_criteria.

Activities carry differentiation — 100.0 (measured, confidence 1.00)

6 of 6 activities carry differentiation.

Activities fit inside a period — 100.0 (measured, confidence 1.00)

6 of 6 activities run within the 40-minute period.

Assessment generation — 97.9

Every item can be marked — 100.0 (measured, confidence 1.00)

24 of 24 items carry an answer or a rubric.

MCQs have one right answer and real distractors — 100.0 (measured, confidence 1.00)

5 of 5 multiple-choice items have three or more options with exactly one marked correct.

Objectives are assessed — 90.0 (measured, confidence 1.00)

9 of 10 learning objectives are reachable from at least one assessment item.

- /knowledge/learning_objectives/0 — LO1_Dates_and_Scope is not assessed

Items span more than recall — 100.0 (measured, confidence 1.00)

5 distinct Bloom level(s) in the bank (analyze, apply, evaluate, remember, understand). Three or more is the point at which a paper distinguishes recall from reasoning.

- /assessments/items — bloom levels: ['analyze', 'apply', 'evaluate', 'remember', 'understand']

Blueprint matches the items actually present — 100.0 (measured, confidence 1.00)

The published blueprint agrees with a recount of the bank.

- /assessments/blueprint/items_by_kind — {'long_answer': 12, 'short_answer': 7, 'mcq': 5}

Total marks equal the sum of the items — 100.0 (measured, confidence 1.00)

Published total is 98; the items sum to 98. They agree.

- /assessments/total_marks — 98 vs 98 summed

Item difficulty is calibrated — n/a (not measurable, confidence 0.00)

Difficulty is a property of a population's responses. No student has answered these items, so there is no p-value or discrimination index to report. Measuring it would need: response data from a cohort — item-response theory needs roughly 100 students per item

Gap analysis — 94.1

Gaps attach to real concepts — 100.0 (measured, confidence 1.00)

7 of 7 predicted gaps name a concept that stage 3 extracted.

Gaps carry a remediation — 100.0 (measured, confidence 1.00)

7 of 7 gaps carry remediation.

Gaps carry a diagnostic question — 100.0 (measured, confidence 1.00)

7 of 7 gaps carry diagnostic_questions.

Diagnostics predict the wrong answer — 100.0 (measured, confidence 1.00)

14 of 14 diagnostic questions state the wrong answer the misconception produces, which is what makes them diagnostic rather than merely difficult.

- /learning_gaps — 14 diagnostic questions

Gaps trace to extracted misconceptions — 0.0 (measured, confidence 1.00)

0 of 7 gaps restate a misconception stage 3 found in the source. The remainder are predicted rather than sourced, which is a legitimate output of this stage.

- /knowledge/misconceptions — 3 extracted misconceptions

Validation — 82.8

No blocking issues remain — 53.3 (measured, confidence 1.00)

Status 'fail' with 15 issue(s), 7 of them blocking.

- /validation/issues — COVERAGE_OBJECTIVE_UNASSESSED: objective 'LO1_Dates_and_Scope' is not covered by any assessment item
- /validation/issues — GROUNDING_UNSUPPORTED_CLAIM: claim at /knowledge/concepts/7 is not supported by its cited source (judge rated this unsupported)
- /validation/issues — GROUNDING_UNSUPPORTED_CLAIM: claim at /knowledge/examples/0 is not supported by its cited source (judge rated this unsupported)
- /validation/issues — GROUNDING_UNSUPPORTED_CLAIM: claim at /knowledge/examples/1 is not supported by its cited source (judge rated this unsupported)

•

/validation/issues — GROUNDING_UNSUPPORTED_CLAIM: claim at /knowledge/applications/1 is not supported by its cited source (judge rated this unsupported)

Coverage report survives a recount — 100.0 (measured, confidence 1.00)

The published coverage numbers match an independent recount of the package.

- /validation/coverage — 25/25 concepts taught, confirmed

Structural checks pass — 100.0 (measured, confidence 1.00)

3 of 3 structural checks are clean.

Claims are supported by the source — 77.8 (measured, confidence 1.00)

Stage 9 scored grounding at 0.78. 14 claim(s) were recorded as unsupported.

- /validation/unsupported_claims — Storming of the Bastille (14 July 1789): The Bastille was attacked on 14 July 1789; though it held few prisoners, its fa
- /validation/unsupported_claims — Constitutional monarchy phase (July 1789 e22 September 1792): A phase in which France had a constitutional monarchy,
- /validation/unsupported_claims — Louis XVI's flight to Varennes (June 1791): The royal family's attempted escape in June 1791 ended with their arrest at
- /validation/unsupported_claims — Reign of Terror (Sept 1793 e22 July 1794): A period of political repression where the Law of Suspects and emergency m
- /validation/unsupported_claims — Thermidorian reaction: The political reaction after Robespierre's fall involving revenge killings in southern France, ex

Publishing — 97.3

Every stage reported its cost — 90.0 (measured, confidence 1.00)

9 of 10 stages recorded a timing entry; missing publishing.

- /provenance/stage_timings — 9 entries totalling 402607 ms

Each model-calling stage names the model that ran it — 100.0 (measured, confidence 1.00)

8 of 8 model-calling stages record which model produced their output. document-intelligence, publishing are excluded because they make no model call, so an absent entry for them is correct rather than missing.

- /generator/models_by_stage — {'educational-classification': 'gpt-5-mini-2025-08-07', 'knowledge-extraction': 'gpt-5-mini-2025-08-07', 'teaching-planner': 'gpt-5-mini-2025-08-07', 'lesson-generation': 'gpt-5-mini-2025-08-07', 'act

Stages explain their choices — 100.0 (measured, confidence 1.00)

8 stage decision(s) recorded with a reason. Four is the point at which the main routing choices — period count, item mix, gap severity, profile — are all accounted for.

- /provenance/decisions — knowledge-extraction: extraction split into 3 passes — 36,184 tokens exceeds the 12,000-token window this route leaves for document text after the schema, system prompt, and the reservation for the answer; sections are packed to that window rather than one call each, so cost follows content volume
- /provenance/decisions — knowledge-extraction: pedagogy pass reads 18 of 80 chunks — objectives and edges are document-level judgements, but the document is 36,184 tokens against a 12,000-token window; scope comes from the full 27-concept inventory, which is sent whole, and the text is narrowed to the passages those concepts cite plus a stride

across the rest, so every claim still has a quotable source

- /provenance/decisions — knowledge-extraction: 2 claims dropped — their quotes did not appear in the chunk they cited; an unverifiable citation is removed here rather than propagating into six later stages
- /provenance/decisions — teaching-planner: 6 periods of 40 minutes — derived from 25 concepts weighted by importance against a 40-minute period

Package is identified and versioned — 100.0 (measured, confidence 1.00)

3 of 3 identity fields present (schema 1.0.0).

- /schema_version — 1.0.0

Cost and token usage are accounted for — 100.0 (measured, confidence 1.00)

\$0.0000 across 164,371 input and 73,508 output tokens.

- /provenance/total_cost_usd — \$0.0000

Learning outcomes — n/a

Teacher satisfaction — n/a (not_measurable, confidence 0.00)

Satisfaction is a rating teachers give. None have used this package. Measuring it would need: a pilot with teachers rating packages after classroom use — 20+ responses before a mean means anything

Student learning effectiveness — n/a (not_measurable, confidence 0.00)

Effectiveness is measured by what students learn, which requires students, a pre-test and a post-test. Measuring it would need: a controlled classroom trial with pre/post assessment against a conventionally-planned unit

Teacher preparation time saved — n/a (not_measurable, confidence 0.00)

Time saved is a difference against how long the same teacher would take unaided. That baseline has not been collected. Measuring it would need: timed preparation of the same unit by teachers with and without the package

Conformance to a published curriculum — n/a (not_measurable, confidence 0.00)

The package aligns to a board's structure, but no authority has certified the mapping. Reporting a conformance score would state agreement nobody granted. Measuring it would need: an official syllabus document parsed to statement level, and a reviewer's sign-off on the mapping